

numpy 1D arrays

go over solutions to homework 2

numpy

- numpy is the main numerical library in Python, basis for many other scientific Python libraries
 - typical usage: `import numpy as np`
 - numpy provides: 1. the `ndarray` object, 2. lots of numerical and array functions
 - arrays are sequences, like lists and tuples, but faster and much more memory efficient
 - ideal for large datasets!
 - unlike lists, can explicitly be multidimensional - useful for e.g. images and movies
 - tradeoff: not as flexible as lists - for efficiency, each entry in an array has to be of the same data type
 - you can have an array of ints, or floats, or strings or booleans, but not a mixture
 - so far, we've seen that there are two main numeric data types: int and float
 - different kinds of integers and floats (see later), each array can contain only one kind
 - like a tuple, array length generally **can't** change, but like a list, its values **can** be changed, so it's "semi-mutable"
- initializing an array
 - explicitly, using a list or a tuple, convert to array:
 - `a = np.array([1, 2, 3])` or `a = np.array((1, 2, 3))`
 - `a = np.arange(10)`
 - very similar to `list(range(10))` , but returns an array instead of a list
 - `a = np.zeros(10)`
 - `a = np.ones(10)`
 - `a = np.random.random(10)` - 10 random numbers uniformly distributed between 0 and 1
 - `a = np.tile([1, 2], 5)`
 - `a.fill(7)` fills the existing array `a` with the number 7
 - array methods (e.g. `a.fill()`) usually operate on the array in-place, while numpy functions (e.g. `np.zeros()`) usually return a new array
 - here's an exception: `b = a.copy()`
 - numpy functions often have array method counterparts (and vice versa)
 - `copy()` and `sort()` are two examples:

```
a = np.random.random(10)
b = a.copy()
c = np.copy(b)
```

 - are `a` , `b` and `c` equal? test with `==` , get a boolean answer for each entry
 - are `a` , `b` and `c` the same objects? test with `is` , get a single bool answer

```
d = a
d.sort() # in-place
e = np.sort(a)
```

 - are `a` , `d` and `e` equal? are they the same objects?

- like other sequences (tuples & lists), get length of array using `len(a)` , but can also get array shape using `a.shape` attribute
 - `shape` returns the length along all dimensions of `a`
 - length of the first dimension is `a.shape[0]` , identical to `len(a)`
 - get num dims with `a.ndim` , multidimensional arrays covered later
- indexing in 1D is the same as for tuples & lists: 0-based, -ve indices count from the end
 - `a[0] = 7` assigns 7 to 1st entry
 - `a[1] = 7` assigns 7 to 2nd entry
 - `a[-1] = 7` assigns 7 to last entry
 - `a[-2] = 7` assigns 7 to 2nd last entry
- slicing in 1D is also the same as for tuples and lists
 - retrieve a slice: the first 5 entries
 - `b = a[0:5]` or `b = a[:5]`
 - assign to a slice: the last 5 entries
 - `a[5:10] = 7` or `a[5:] = 7`
 - assign to a slice: all entries
 - `a[:] = 8` , same as `a.fill(8)`
 - what happens if you go `a = 8` ?
- arrays also have "fancy" indexing:
 - allow you to ask for multiple values from an array at once
 - two types: integer & boolean fancy indexing
 - both are kind of hybrid between normal indexing and slicing
 - benefit over slicing is that you can specify any sequence of indices, not just evenly spaced ones
 - you can even specify the same index multiple times
 - integer fancy indexing
 - ```
a = np.random.random(10) # init an array of random data
i = [3, 7, 5, 2, 7] # create a list of indices
vals = a[i] # index into array using integer fancy indexing
a[i] = -1 # assign -1 at multiple locations using integer fancy indexing
```
  - can ask for array values in arbitrary order
    - can ask for the same value repeatedly
    - can't do this with lists: try `list(range(10))[i]`
  - boolean fancy indexing
    - ask some question of values of the array, get an answer back of boolean values of same length as original array
    - `i = a > 5` returns an array of booleans, which can be used for indexing
    - `a[a > 5]` or `a[i]` returns only those entries in `a` that are `> 5`
    - i.e., where `i` is True, return the value in `a` at that index
    - what if you have another array `b` that is of different length? can you also index into it with the above `i` ? no!
    - again, can't do this with lists: try `l[i]`
- **vectorized** math operators ( `+` , `-` , `*` , `/` , `**` ) and comparitors ( `==` , `>` , `<` , `!=` )

- vectorized: work on all values of an array at the same time
- `a = np.array([1, 2, 3])`
- arrays & scalars
  - `a + 1` returns a new array with 1 added to all the values in `a`
  - `a += 1` increments `a` in-place by 1, doesn't return anything
  - `a -= 1` decrements `a` in-place by 1, doesn't return anything
- `b = np.array([4, 5, 6])`
- `a + b` returns another array each of whose values are the sum of the corresponding two values in `a` and `b`
  - in comparison, what does `+` do for strings and lists?
  - use `np.concatenate([a, b])` or `np.concatenate((a, b))` to combine arrays
- what happens if you try to do one of the above vectorized operations on two arrays of different length?

### array exercises:

1. Use a for loop to build a list of 3 arrays, each array of length 5, initialized to zeros
2. Find the average difference between the following two arrays: `a = np.array([10, 20, 30, 40, 50])`, `b = np.array([5, 10, 15, 20, 25])`
3. Write a function called `rms()` that calculates the RMS (root mean square) of an input sequence (array, list, tuple). RMS is the the square root of the mean of the square of a signal.
4. Use your `rms()` function to calculate the RMS of the difference between the two arrays

<go over solutions>

5. Create an array `c` of 10 random numbers between 0 and 10.
6. Create an array `d` that consists of the values in `c` greater than 7.
7. Create an array `e` of 10 random numbers between -5 and 5.
8. Create an array `f` that consists of values in `e` greater than 1 and less than -1.
9. Create an array `g` that has all the values of both `d` and `f`. How long is it?
10. Create an array `h` that has only the 3rd, 8th and 11th entries in `g`

- memory

- what's system memory (RAM)? computer's working memory
- what's a byte? 8 bits
- what's a bit? a binary digit, can be a 0 or 1
- different numeric values are expressed using different combinations of bits
  - 1 byte, 8 bits allow for  $2^{**8} = 256$  different numeric values to be expressed
  - `00000000`, `00000001`, `00000010`, `00000011` ... == 0, 1, 2, 3, ...
- how much memory does my array use?
  - `a.nbytes`
  - memory use depends on the number of elements in the array, times the size of each element
  - element size depends on the data type (dtype) of the array - `a.dtype`
  - `a.nbytes == len(a) * a.dtype.itemsize` for 1D arrays

- array data type (dtype)

- a common set of numeric data types are used across programming languages, super important!
- integers: signed and unsigned
  - signed integers are symmetric around 0, unsigned integers are always  $\geq 0$ 
    - if  $n$  is the number of unique integers that can be represented by an integer data type:
      - signed integers range from  $-n/2$  to  $n/2-1$
      - unsigned integers range from  $0$  to  $n-1$
      - $n = 2^{*}nbits$
      - so, the bigger the integer data type (in bits or bytes), the more integer numbers it can represent
  - `np.int8`, `np.int16`, `np.int32`, `np.int64` - 1, 2, 4 and 8 byte signed
  - `np.uint8`, `np.uint16`, `np.uint32`, `np.uint64` - 1, 2, 4 and 8 byte **unsigned**
  - can easily calculate max/min values of each int dtype, or use `np.iinfo()`, e.g. `np.iinfo(np.int8).max`
  - init arrays to the desired data type by using the `dtype` kwarg:
    - `a = np.zeros(10, dtype=np.uint8)`
    - `b = np.zeros(10, dtype=np.int64)`
  - integer overflow:
    - `a[:] = 255` is fine, but `a[:] = 256` and `a[:] = -1` both wrap overflow (wrap around)
    - `b[:] = 127` is fine, but `a[:] = -128` isn't
    - `b[:] = -128` is fine, but `b[:] = -129` isn't
  - when to use signed or unsigned? if in doubt, use signed!
- floats - always signed, and made of "mantissa +  $10^{\text{exponent}}$ "
  - bigger floats have greater precision
  - `np.float16`, `np.float32`, `np.float64` - 2, 4 and 8 bytes floats
- by default, arrays init to the biggest dtypes, either `np.float64` or `np.int64`:
  - `a = np.array([1, 2, 3])`, `a.dtype` -> `int64`
  - `b = np.array([1.1, 2.2, 3.3])`, `b.dtype` -> `float64`
- init arrays to the desired data type by using the `dtype` kwarg:
  - `a = np.zeros(10, dtype=np.int8)`
  - `b = np.zeros(10, dtype=np.int64)`
  - `c = np.zeros(10, dtype=np.float64)`
  - calculate how much memory do `a`, `b` and `c` take, then check it using `.nbytes`
- how much memory would `a = np.zeros(20000000000, dtype=np.uint8)` use? what would happen if I tried this on my 16 GB laptop? `MemoryError`
- can convert from one dtype to another by using the dtype as a function:
  - `a = np.array([1, 2, 3])`
  - `np.float64(a)` converts `a` to `float64` dtype
    - similar to basic Python: `float(val)`
  - `a = np.array([1.1, 2.2, 3.3])`
  - `np.int64(a)` converts `a` to `int64` dtype, but it truncates!
    - this is similar to basic Python: `int(val)`

- use `np.int64(np.round(a))` to round to the nearest integer instead
  - check array data type with `a.dtype`
- usually only need to worry about int vs float dtype, stick to the defaults `int64` and `float64`
  - only consider going down to smaller dtypes if you have lots of data and not enough memory on your machine
- take care converting between dtypes!
  - especially from larger ones to smaller ones, and from floats to ints
  - a number that can be represented in one data type might not be possible to represent in another
  - dramatic example: Ariane 5 1996 failure
    - code adapted from Ariane 4 tried to convert a large float64 to int16, resulted in integer overflow, caused computer to think it was suddenly way off course, tried to correct by rapidly changing direction, high G-forces caused it to start to disintegrate, which triggered self-destruct. Cost: \$370M

### **quickie numeric data type exercises:**

1. Create a tuple or a list with the following entries: `3, 5, 1.7, -2.7, 1e2, -50`. Now convert it to an array. What happens?
2. You have integer data whose values span -100 to 100. Normally, you would use an `int64` array to store this data, except the dataset is huge (1 billion entries) and your laptop has a measly 4 GB of RAM. How much memory would you need to store your data using the default `int64` dtype? What would be the optimal dtype to use to minimize memory use by your dataset? Will it fit into your 4 GB of RAM?